

Supplementary material for Lesson 10

1 Conjugate Posterior for the Normal Mean

Here we derive the update formula from Lesson 10.1 where the likelihood is normal and we use a conjugate normal prior on the mean. Specifically, the model is $x_1, x_2, \dots, x_n \stackrel{\text{iid}}{\sim} N(\mu, \sigma_0^2)$, $\mu \sim N(m_0, s_0^2)$ with σ_0^2 , m_0 , and s_0^2 known. First consider the case in which we have only one data point x . The posterior is then

$$\begin{aligned}
f(\mu|x) &= \frac{f(x|\mu)f(\mu)}{\int_{-\infty}^{\infty} f(x|\mu)f(\mu)d\mu} \\
&\propto f(x|\mu)f(\mu) \\
&= \frac{1}{\sqrt{2\pi\sigma_0^2}} \exp\left\{-\frac{1}{2\sigma_0^2}(x-\mu)^2\right\} \frac{1}{\sqrt{2\pi s_0^2}} \exp\left\{-\frac{1}{2s_0^2}(\mu-m_0)^2\right\} \\
&\propto \exp\left\{-\frac{1}{2\sigma_0^2}(x-\mu)^2\right\} \exp\left\{-\frac{1}{2s_0^2}(\mu-m_0)^2\right\} \\
&= \exp\left\{-\frac{1}{2\sigma_0^2}(x-\mu)^2 - \frac{1}{2s_0^2}(\mu-m_0)^2\right\} \\
&= \exp\left\{-\frac{1}{2}\left[\frac{1}{\sigma_0^2}(x^2 - 2x\mu + \mu^2) + \frac{1}{s_0^2}(\mu^2 - 2m_0\mu + m_0^2)\right]\right\} \\
&= \exp\left\{-\frac{1}{2}\left[\frac{\mu^2}{\sigma_0^2} + \frac{\mu^2}{s_0^2} + \frac{-2x}{\sigma_0^2}\mu + \frac{-2m_0}{s_0^2}\mu + \frac{x^2}{\sigma_0^2} + \frac{m_0^2}{s_0^2}\right]\right\} \\
&= \exp\left\{-\frac{1}{2}\left[\left(\frac{1}{s_0^2} + \frac{1}{\sigma_0^2}\right)\mu^2 - 2\left(\frac{x}{\sigma_0^2} + \frac{m_0}{s_0^2}\right)\mu + \frac{x^2}{\sigma_0^2} + \frac{m_0^2}{s_0^2}\right]\right\} \\
&= \exp\left\{-\frac{1}{2}\left[\left(\frac{1}{s_0^2} + \frac{1}{\sigma_0^2}\right)\mu^2 - 2\left(\frac{x}{\sigma_0^2} + \frac{m_0}{s_0^2}\right)\mu\right]\right\} \cdot \exp\left\{-\frac{1}{2}\left[\frac{x^2}{\sigma_0^2} + \frac{m_0^2}{s_0^2}\right]\right\} \\
&\propto \exp\left\{-\frac{1}{2}\left[\left(\frac{1}{s_0^2} + \frac{1}{\sigma_0^2}\right)\mu^2 - 2\left(\frac{x}{\sigma_0^2} + \frac{m_0}{s_0^2}\right)\mu\right]\right\} \\
&= \exp\left\{-\frac{1}{2}\left[\frac{1}{s_0^2}\mu^2 - 2\left(\frac{x}{\sigma_0^2} + \frac{m_0}{s_0^2}\right)\mu\right]\right\} \quad \text{where } s_1^2 = \left(\frac{1}{s_0^2} + \frac{1}{\sigma_0^2}\right)^{-1} \\
&= \exp\left\{-\frac{1}{2}\left[\frac{1}{s_1^2}\mu^2 - 2\frac{s_1^2}{s_0^2}\left(\frac{x}{\sigma_0^2} + \frac{m_0}{s_0^2}\right)\mu\right]\right\} \\
&= \exp\left\{-\frac{1}{2s_1^2}[\mu^2 - 2m_1\mu]\right\} \quad \text{where } m_1 = s_1^2 \left(\frac{m_0}{s_0^2} + \frac{x}{\sigma_0^2}\right).
\end{aligned}$$

The next step is to “complete the square” in the exponent:

$$\begin{aligned}
f(\mu|x) &\propto \exp \left\{ -\frac{1}{2s_1^2} [\mu^2 - 2m_1\mu] \right\} \\
&= \exp \left\{ -\frac{1}{2s_1^2} [\mu^2 - 2m_1\mu + m_1^2 - m_1^2] \right\} \\
&= \exp \left\{ -\frac{1}{2s_1^2} [\mu^2 - 2m_1\mu + m_1^2] \right\} \cdot \exp \left\{ \frac{m_1^2}{2s_1^2} \right\} \\
&\propto \exp \left\{ -\frac{1}{2s_1^2} [\mu^2 - 2m_1\mu + m_1^2] \right\} \\
&= \exp \left\{ -\frac{1}{2s_1^2} (\mu - m_1)^2 \right\}
\end{aligned}$$

which, except for a normalizing constant not involving μ , is the PDF of a normal distribution with mean m_1 and variance s_1^2 .

The final step is to extend this result to accommodate n independent data points. The likelihood in this case is

$$\begin{aligned}
f(\mathbf{x}|\mu) &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma_0^2}} \exp \left\{ -\frac{1}{2\sigma_0^2} (x_i - \mu)^2 \right\} \\
&= (2\pi\sigma_0^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma_0^2} \sum_{i=1}^n (x_i - \mu)^2 \right\} \\
&\propto \exp \left\{ -\frac{1}{2\sigma_0^2} \left[\sum_{i=1}^n x_i^2 - 2\mu \sum_{i=1}^n x_i + n\mu^2 \right] \right\} \\
&\propto \exp \left\{ -\frac{1}{2\sigma_0^2} [-2n\bar{x}\mu + n\mu^2] \right\}.
\end{aligned}$$

We can repeat the steps above *or* notice that the data contribute only through the sample mean \bar{x} (and n which we assume is known). This means that \bar{x} is a “sufficient statistic” for μ , allowing us to use the distribution of \bar{x} as the likelihood (analogous to using a binomial likelihood in place of a sequence of Bernoullis). The model then becomes

$$\bar{x} \mid \mu \sim N\left(\mu, \frac{\sigma_0^2}{n}\right), \quad \mu \sim N(m_0, s_0^2).$$

We now apply our result derived above, replacing x with \bar{x} and σ_0^2 with σ_0^2/n . This yields the update equation presented in Lesson 10.1.

2 Marginal Distribution of Normal Mean in Conjugate Model

Consider again the model $x|\mu \sim N(\mu, \sigma_0^2)$, $\mu \sim N(m_0, s_0^2)$ with σ_0^2 known. Here we derive the marginal distribution for data given by $\int_{-\infty}^{\infty} f(x|\mu)f(\mu)d\mu$. This is the prior predictive distribution for a new data point x^* .

To do so, re-write the model in an equivalent, but more convenient form: $x = \mu + \epsilon$ where $\epsilon \sim N(0, \sigma_0^2)$ and $\mu = m_0 + \eta$ where $\eta \sim N(0, s_0^2)$, with ϵ and η independent. Now substitute μ into the first equation to get $x = m_0 + \eta + \epsilon$. Recall that adding two normal random variables results in another normal random variable, so x is normal with $E(x) = E(m_0 + \eta + \epsilon) = m_0 + E(\eta) + E(\epsilon) = m_0 + 0 + 0$ and $Var(x) = Var(m_0 + \eta + \epsilon) = Var(m_0) + Var(\eta) + Var(\epsilon) = 0 + s_0^2 + \sigma_0^2$ (note that we can add variances because of the independence of η and ϵ). Therefore the marginal distribution for x is normal with mean m_0 and variance $s_0^2 + \sigma_0^2$. The posterior predictive distribution is the same, but with m_0 and s_0^2 replaced by the posterior updates given in Lesson 10.1.

3 Inverse-Gamma Distribution

The inverse-gamma distribution is the conjugate prior for σ^2 in the normal likelihood with known mean. It is also the marginal prior/posterior for σ^2 in the model of Lesson 10.2.

As the name implies, the inverse-gamma distribution is related to the gamma distribution. If $X \sim \text{Gamma}(\alpha, \beta)$, then the random variable $Y = 1/X \sim \text{Inverse-Gamma}(\alpha, \beta)$ where

$$f(y) = \frac{\beta^\alpha}{\Gamma(\alpha)} y^{-(\alpha+1)} \exp\left(-\frac{\beta}{y}\right) I_{\{y>0\}}$$

$$E(Y) = \frac{\beta}{\alpha - 1} \quad \text{for } \alpha > 1.$$

The relationship between gamma and inverse-gamma suggest a simple method for simulating draws from the inverse-gamma distribution. First draw X from the $\text{Gamma}(\alpha, \beta)$ distribution and take $Y = 1/X$, which corresponds to a draw from the $\text{Inverse-Gamma}(\alpha, \beta)$.

4 Marginal Posterior Distribution for the Normal Mean when the Variance is Unknown

If we are not interested in inference for an unknown σ^2 , we can integrate it out of the joint posterior in Lesson 10.2. This results in a t-distributed marginal posterior as noted at the end of the lesson. This t distribution has $\nu = 2\alpha + n$ degrees of freedom and two additional parameters, a scale γ and a location m^* given by

$$m^* = \frac{n\bar{x} + w\bar{m}}{n + w} \quad (\text{the mean of the conditional posterior for } \mu)$$

$$\gamma = \sqrt{\frac{\beta + \frac{n-1}{2}s^2 + \frac{wn}{2(w+n)}(\bar{x} - m)^2}{(n+w)(\alpha+n/2)}} \quad (\text{modified scale of the updated inverse-gamma for } \sigma^2)$$

where $s^2 = \sum(x_i - \bar{x})^2/(n - 1)$, the sample variance.

This t distribution can be used to create a credible interval for μ by multiplying the appropriate quantiles of the standard t distribution by the scale γ and adding the location m^* .

Example: Suppose we have normal data with unknown mean μ and variance σ^2 . We use the model from Lesson 10.2 with $m = 0$, $w = 0.1$, $\alpha = 3/2$, and $\beta = 1$. The data are $n = 20$ independent observations with $\bar{x} = 1.2$ and $s^2 = 0.7$. Then we have

$$\begin{aligned} \sigma^2 | \mathbf{x} &\sim \text{Inverse-Gamma}(11.5, 7.72) \\ \mu | \sigma^2, \mathbf{x} &\sim N\left(1.19, \frac{\sigma^2}{20.1}\right) \\ m^* &= 1.19 \\ \gamma &= 0.183 \end{aligned}$$

and $\mu | \mathbf{x}$ is distributed t with 23 degrees of freedom, location 1.19 and scale 0.183. To produce a 95% equal-tailed credible interval for μ , we first need the 0.025 and 0.975 quantiles of the standard t distribution with 23 degrees of freedom. These are -2.07 and 2.07 . The 95% credible interval is then $m^* \pm \gamma(2.07) = 1.19 \pm 0.183(2.07) = 1.19 \pm 0.38$.